

Computação Paralela

Mest. Engenharia Computacional
Mest. Int. Engenharia Computacional

Ano letivo 2022/2023

Rui Costa, Nuno Lau

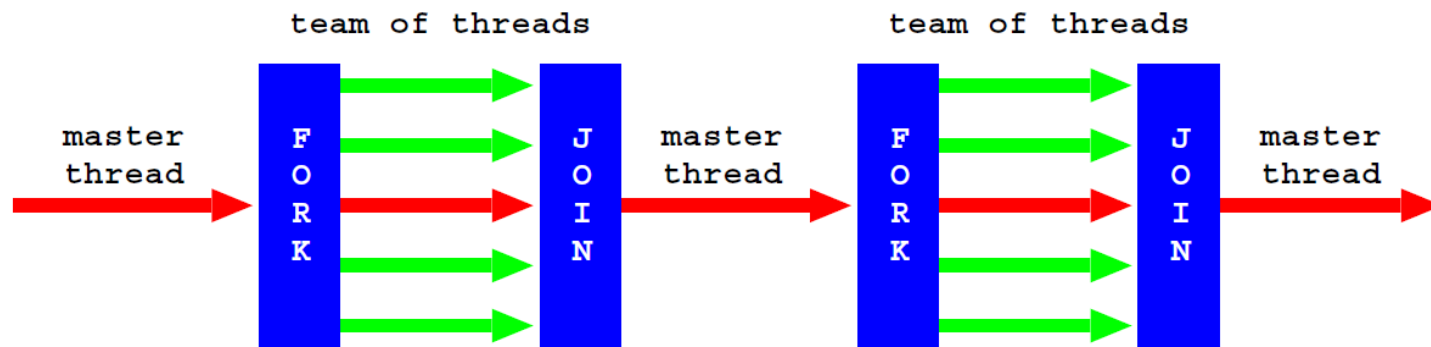
- **Open specifications for Multi Processing** *via collaborative work between interested parties from the hardware and software industry, government and academia*

Acknowledgement: This lecture is based on “A ‘Hands-on’ Introduction to OpenMP”, Tim Mattson, Intel Corp., timothy.g.mattson@intel.com

- **Shared Memory Programming Model**
- Cooperation of several hardware and software companies (AMD, Intel, arm, Fujitsu, IBM, HP, NASA, NEC, NVIDIA, Siemens, SUSE, ...)
- Parallel Programming API for multiprocessor / multicore architectures
- Languages: C/C++ or Fortran
- OS: Unix/Linux or Windows
- **OpenMP is a specification not an implementation!**

OpenMP fork-join execution

- Program initiates with a single (master) thread
- Executes sequentially until parallel region defined by OpenMP constructor, and then:
 - Master thread forks “team of threads”
 - Parallel region code is executed concurrently by all threads (including master thread)
 - There is an implicit barrier at the end of parallel region
 - “team of threads” terminates and master thread continues sequentially



From CP@FCUP

- OpenMP has constructs to enable mutual exclusion among threads
- 3 different ways:
 - **omp critical**
 - Defines critical region
 - Can be used for any code
 - **omp atomic**
 - Mutual exclusion for atomic variable updates
 - Only for simple memory updates
 - $x \text{ binop} = \text{expr}$, $x++$; $++x$; $x--$; $--x$
 - Explicit use of locks

- **omp critical**

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    #pragma omp parallel  
    {  
        int private_r = 0;  
        #pragma omp for  
        for (int i = 0; i < n; i++)  
            private_r += u[i] * v[i];  
        #pragma omp critical  
        r += private_r;  
    }  
    return r;  
}
```

- **omp atomic**

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    #pragma omp parallel  
    {  
        int private_r = 0;  
        #pragma omp for  
        for (int i = 0; i < n; i++)  
            private_r += u[i] * v[i];  
        #pragma omp atomic  
        r += private_r;  
    }  
    return r;  
}
```

- Locks

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    omp_lock_t lock;  
    omp_init_lock(&lock); // initialize lock  
    #pragma omp parallel  
    {  
        int private_r = 0;  
        #pragma omp for  
        for (int i = 0; i < n; i++)  
            private_r += u[i] * v[i];  
        omp_set_lock(&lock); // acquire lock  
        r += private_r;  
        omp_unset_lock(&lock); // release lock  
    }  
    omp_destroy_lock(&lock); // destroy lock  
    return r;  
}
```


- **omp master**
 - Executed only by master thread
 - Ignored by others; no synchronization
- **omp single**
 - Executed by a single thread
 - Implicit barrier for all threads
 - Except `nowait` clause is used
- **omp barrier**
 - Explicit barrier

- **omp master**
 - Executed only by master thread
 - Ignored by others; no synchronization
- **omp single**
 - Executed by a single thread
 - Implicit barrier for all threads
 - Except `nowait` clause is used
- **omp barrier**
 - Explicit barrier

OpenMP `schedule` clause

- **`schedule (static [, chunk])`**
 - Each thread has fixed number of iterations
- **`schedule (dynamic [, chunk])`**
 - Each thread grabs “chunk” iterations off a queue until all iterations have been handled
- **`schedule (guided [, chunk])`**
 - Threads dynamically grab blocks of iterations
 - The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds.
- **`schedule (runtime)`**
 - Schedule and chunk size taken from the `OMP_SCHEDULE` environment variable (or the runtime library).
- **`schedule (auto)`**
 - Schedule is left up to the runtime to choose (does not have to be any of the above).

nowait ordered lastprivate clauses

- **nowait**
 - threads do not need to synchronize at the end of cycle
- **ordered**
 - Allows `#pragma ordered` blocks that must execute in the cycle iteration order
- **lastprivate(list)**
 - Variables in list are private
 - Their last value is transported to after the cycle

nowait ordered lastprivate clauses

```
#pragma omp for nowait
for (int i = 0; i < n; i++) a[i] = x*i;
#pragma omp for
for (int i = 0; i < n; i++) b[i] = i*i;
#pragma omp for
for (int i = 0; i < n; i++) c[i] = a[i] + b[i];
#pragma omp for ordered lastprivate(x)
for (int i = 1; i < n; i++)
#pragma omp ordered
{
    c[i] += c[i-1]; x = c[i];
    printf("c[%d] = %d\n", c[i]);
}
```

- threads may go to 2nd cycle even before all threads terminated 1st cycle (**nowait**)
- In the last cycle, each thread only executes after previous iterations are concluded (**ordered**)
- **x** gets the value of c[n-1] (**lastprivate**)

OpenMP `omp sections` directive

```
#pragma omp sections [clause, ...]
{
    #pragma omp section
    { ... }
    ...
    #pragma omp section
    { ... }
}
```

- The `omp sections` directive defines a set of code sections that can perform concurrently, allowing functional parallelism between sections.
- Each section is identified by an `omp section` directive and is performed by just one thread.
- All threads, including those that are not involved in any section, synchronize at the end of the `omp sections` region, unless the `nowait` clause is specified.

OpenMP `omp sections` directive

```
x = f1 ();  
a = f2 (x) ;  
b = f3 (x, 1) ;  
c = f4 (x, a) ;  
d = f5 (a, b) ;  
e = f6 (d) ;
```

- Assuming that **f1** to **f6** can execute concurrently correct, how to parallelize the code?
- Dependencies allow for the following processing order:
 - x** = **f1** () for just one thread.
 - a** = **f2** (**x**) and **b** = **f3** (**x**, 1) for 2 threads in parallel.
 - c** = **f4** (**x**, **a**) and **d** = **f5** (**a**, **b**) for 2 threads in parallel.
 - e** = **f6** (**d**) for only one thread.

OpenMP omp sections directive

```
#pragma omp parallel
{
    #pragma omp single
    x = f1();
    #pragma omp sections
    {
        #pragma omp section
        a = f2(x);
        #pragma omp section
        b = f3(x,1);
    }
    #pragma omp sections
    {
        #pragma omp section
        c = f4(x, a);
        #pragma omp section
        d = f5(a,b);
    }
    #pragma omp single
    e = f6(d);
}
```


- A task has
 - Code to execute
 - Data environment (it owns its data)
 - An assigned thread that executes the code and used the data
- Two activities: packaging and execution
 - Each encountering thread packages a new instance of a task (code and data)
 - Some thread in the team executes the task at some later time

- Tasks have been fully integrated into OpenMP
- Key concept: OpenMP has always had tasks, we just never called them that.
 - Thread encountering parallel construct packages up a set of implicit tasks, one per thread.
 - Team of threads is created.
 - Each thread in team is assigned to one of the tasks (and tied to it).
 - Barrier holds original master thread until all implicit tasks are finished.
- We have simply added a way to create a task explicitly for the team to execute.
- Every part of an OpenMP program is part of one task or another!

OpenMP task example

- **omp task** used to process elements of a linked list:

```
#pragma omp parallel
{
    #pragma omp single private(p)
    {
        p = listhead ;
        while (p) {
            #pragma omp task
            process (p)
            p = next(p) ;
        }
    }
}
```