

Computação Paralela

Mest. Engenharia Computacional
Mest. Int. Engenharia Computacional


Ano letivo 2022/2023

Rui Costa, Nuno Lau


- Parallel general purpose computation model
- Introduced by NVIDIA in 2006
- Allows using the GPU for the execution of general purpose applications

Acknowledgement: This lecture is based on “CUDA C++ Programming Guide”, v11.6, NVIDIA and older versions


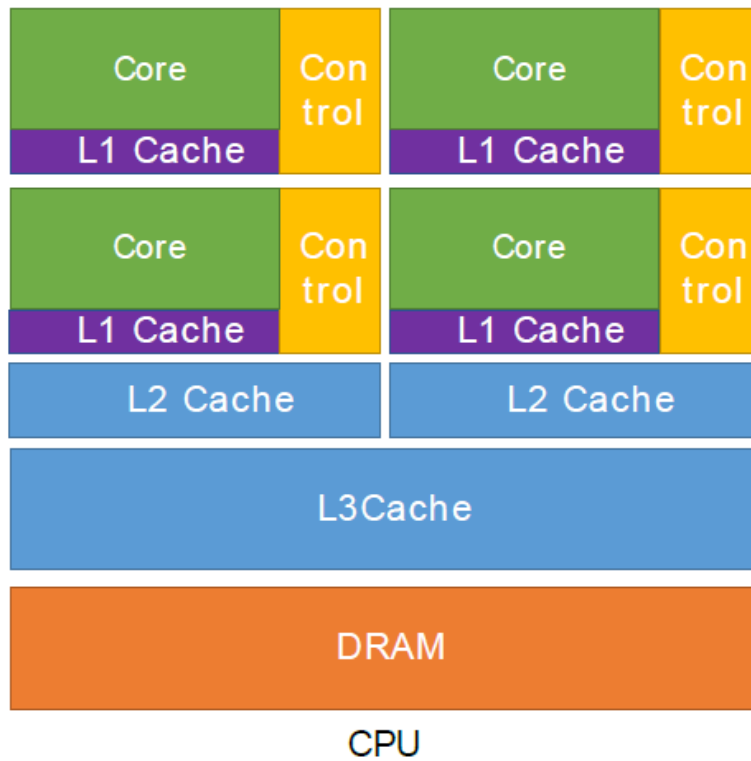
Theoretical GFLOP/s



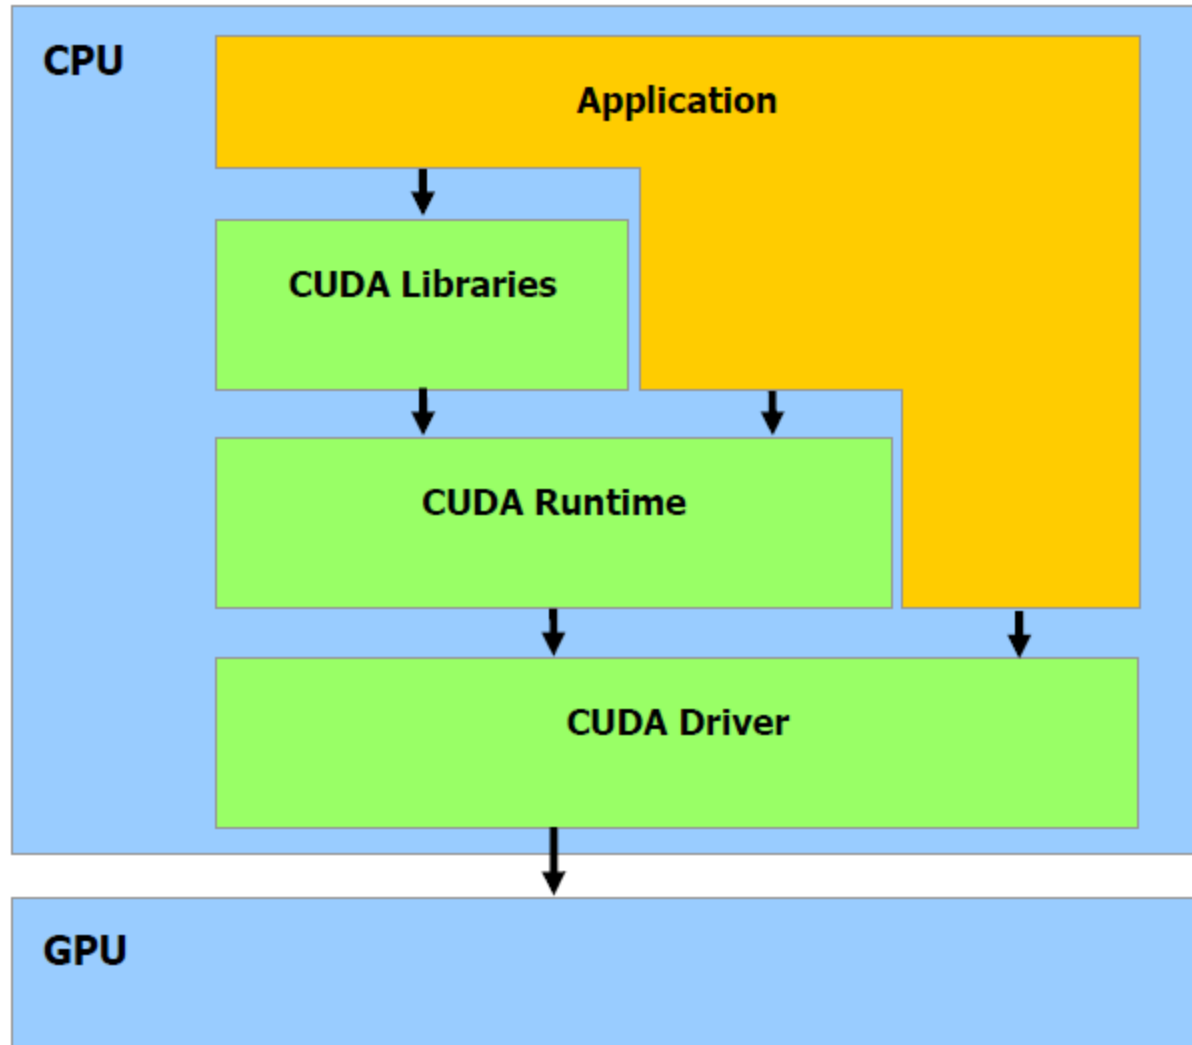
Theoretical GB/s








CUDA – CPU and GPU Architecture




CUDA Architecture



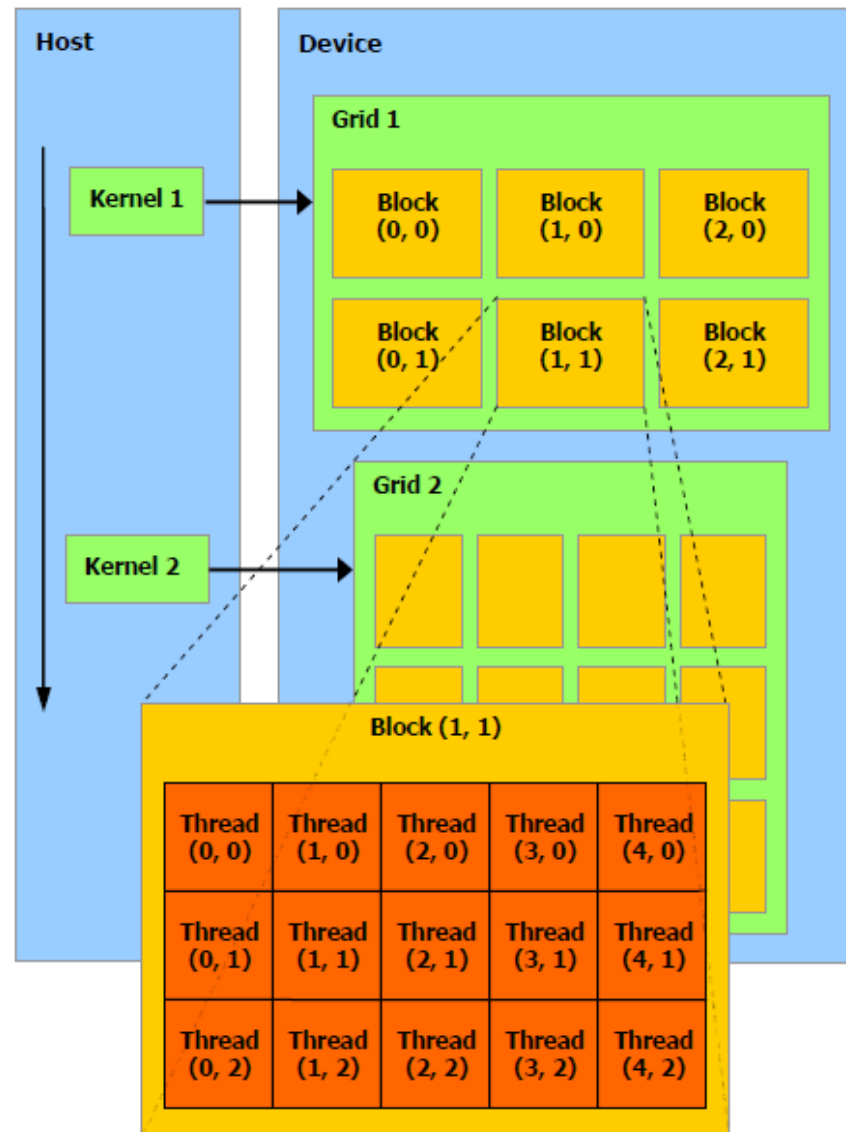
CUDA Computing Applications

GPU Computing Applications						
Libraries and Middleware						
cuDNN TensorRT	cuFFT cuBLAS cuRAND cuSPARSE	CULA MAGMA	Thrust NPP	VSIPL SVM OpenCurrent	PhysX OptiX iRay	MATLAB Mathematica
Programming Languages						
C	C++	Fortran	Java Python Wrappers	DirectCompute	Directives (e.g. OpenACC)	
			CUDA-Enabled NVIDIA GPUs			
NVIDIA Ampere Architecture (compute capabilities 8.x)					Tesla A Series	
NVIDIA Turing Architecture (compute capabilities 7.x)		GeForce 2000 Series	Quadro RTX Series	Tesla T Series		
NVIDIA Volta Architecture (compute capabilities 7.x)	DRIVE/JETSON AGX Xavier		Quadro GV Series	Tesla V Series		
NVIDIA Pascal Architecture (compute capabilities 6.x)	Tegra X2	GeForce 1000 Series	Quadro P Series	Tesla P Series		
	 Embedded	 Consumer Desktop/Laptop	 Professional Workstation	 Data Center		

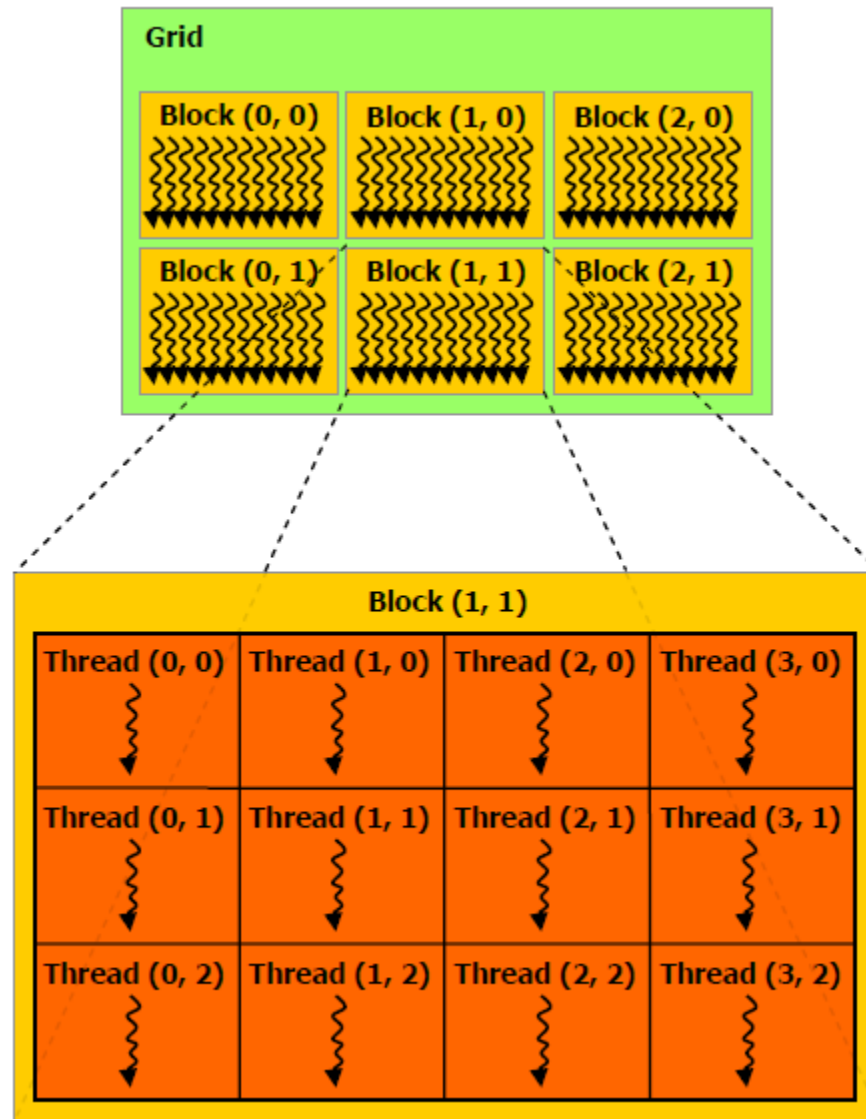
CUDA Automatic Scalability



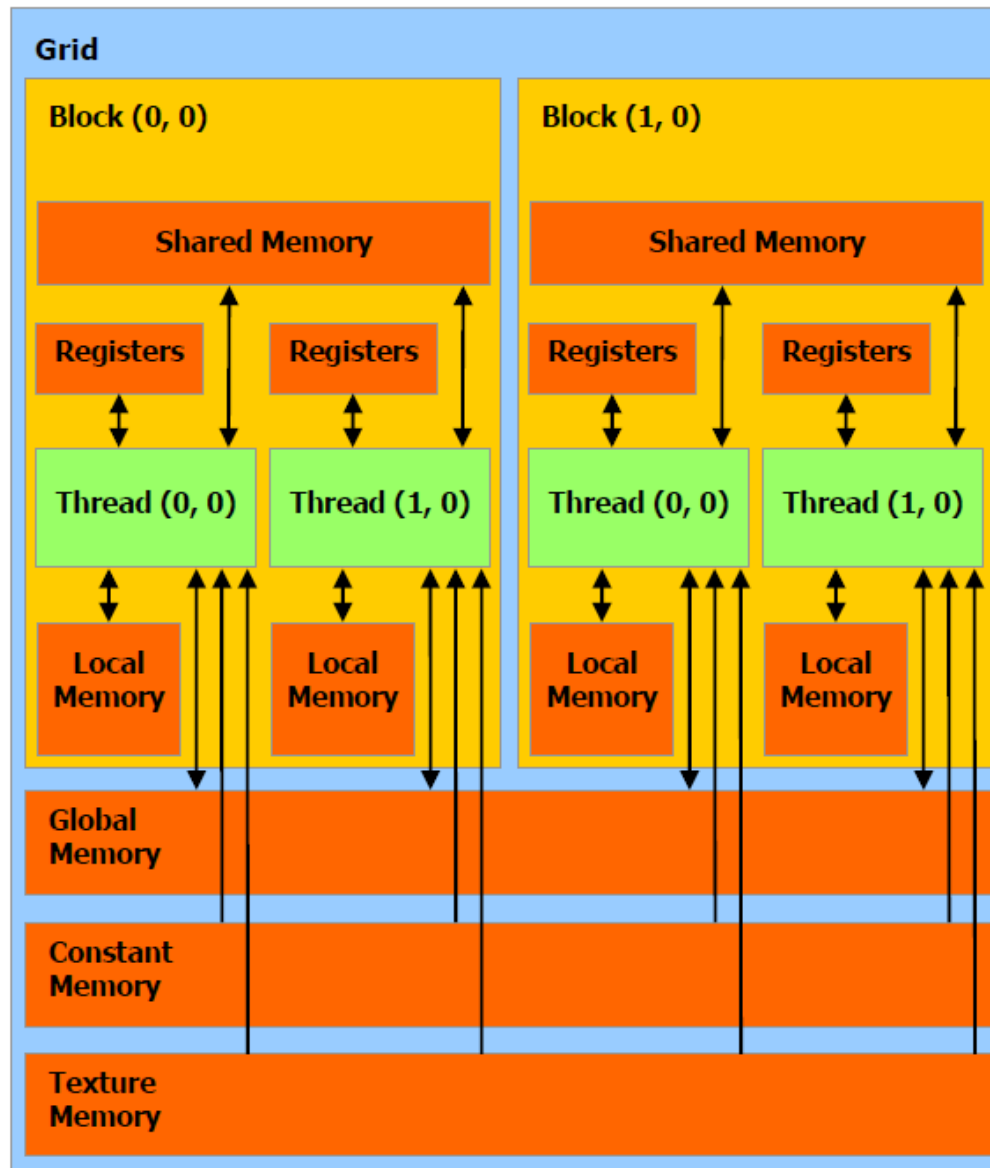
CUDA Kernel, Grid, Block, Thread



CUDA Grid of Threads



CUDA memory



CUDA memory

Thread




Per-thread local
memory

Thread Block



Per-block shared
memory

CUDA memory



Compute Capability 1.x

- Global memory (read and write)
 - Slow and uncached
 - Requires sequential and aligned 16 byte reads and writes to be fast
- Texture memory (read only)
 - Cache optimized for 2D spatial access pattern
- Constant memory
 - This is where constants and kernel arguments are stored
 - Slow, but with small cache
- Shared memory (16 kb per MP)
 - Fast, but take care of bank conflicts
 - Can be used to exchange data between threads in a block
- Local memory (used for data that does not fit into registers)
 - Slow and uncached
- Registers
 - Fastest, scope is thread local

Compute Capability 2.x

- Global memory (read and write)
 - Slow, but now with cache
- Texture memory (read only)
 - Cache optimized for 2D spatial access pattern
- Constant memory
 - Slow, but with cache (8 kb)
- Shared memory
 - Fast, but slightly different rules for bank conflicts now
- Local memory
 - Slow, but now with cache
- Registers (32768 32-bit registers per MP)

Compute Capability 3.x

- Global memory (read and write)
 - Generally cached in L2, but not in L1
- Constant memory
 - Cache shared by all functional units
- Shared memory
 - 32 banks with two addressing modes
- Local memory
 - Cached in L1

Compute Capability 5.x

- Global memory
 - Identical to cp 3.x
- Shared memory
 - 32 banks organized such that successive 32-bit words map to successive banks

Compute Capability 6.x

Compute Capability 7.x

Compute Capability 8.x

CUDA Heterogeneous Programming

C Program
Sequential
Execution

Serial code

Parallel kernel
Kernel0<<<>>>()

Serial code

Parallel kernel
Kernel1<<<>>>()

Host

Device

Grid 0

Block (0, 0)

Block (1, 0)

Block (2, 0)

Block (0, 1)

Block (1, 1)

Block (2, 1)

Host

Device

Grid 1

Block (0, 0)

Block (1, 0)

Block (0, 1)

Block (1, 1)

Block (0, 2)

Block (1, 2)

- Kernel invocation

kernel<<< Dg, Db, Ns, S >>>

- **Dg** is the grid dimension (dim3 type)
- **Db** is the block dimension (dim3 type)
- **Ns** is the number of shared memory bytes
- **S** is the stream

- Block
 - `threadIdx.x`, `threadIdx.y`, `threadIdx.z`
 - Identify the thread in the block
 - `blockDim.x`, `blockDim.y`, `blockDim.z`
 - `threadID = x+y*Dx+z*Dx*Dy`
 - Threads are executed in *warps*
 - 32 threads of the same block with consecutive ids
 - Blocks should have more than 32 threads

```
// Kernel definition
__global__ void VecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x;
    C[i] = A[i] + B[i];
}

int main()
{
    ...
    // Kernel invocation with N threads
    VecAdd<<<1, N>>>(A, B, C);
}
```

- Grid
 - `blockIdx.x, blockIdx.y, blockIdx.z`
 - Identify the block in the grid
 - `gridDim.x, gridDim.y, gridDim.z`

- **__global__**
 - Used to define kernel functions
 - Called by CPU; Run in GPU
- **__device__**
 - Used to define device functions
 - Called by GPU; Run in GPU
- **__shared__**
 - Used to declare variables that reside in shared memory

CUDA main()

```
int main(void)
{
    float A[SIZE], B[SIZE], C[SIZE];
    float *devPtrA; float *devPtrB; float *devPtrC;
    int memsize = SIZE * sizeof(float);

    // Initialize arrays
    srand (time(NULL));
    for(int i=0; i < SIZE; i++) {
        A[i]=rand() % 100; B[i]=rand() % 100;
    }

    cudaSetDevice(0);    // Select GPU device (can be 0 to 1)

    // Allocate device memory for A, B and C arrays
    cudaMalloc((void**)&devPtrA, memsize);
    cudaMalloc((void**)&devPtrB, memsize);
    cudaMalloc((void**)&devPtrC, memsize);

    // Copy data (data to process) from host to device (from CPU to GPU)
    cudaMemcpy(devPtrA, A, memsize, cudaMemcpyHostToDevice);
    cudaMemcpy(devPtrB, B, memsize, cudaMemcpyHostToDevice);

    // Execute the Kernel
    arrAdd <<<1, SIZE>>> (devPtrA, devPtrB, devPtrC); // launch 1 block with SIZE threads

    // Copy data from device (results) back to host
    cudaMemcpy(C, devPtrC, memsize, cudaMemcpyDeviceToHost);

    // Show results
    printf("      A      B      C\n");
    for (int i=0; i < SIZE; i++) {
        printf("%2d: %4.1f + %4.1f = %5.1f\n", i, A[i], B[i], C[i]);
    }

    // Free device memory
    cudaFree(devPtrA); cudaFree(devPtrB); cudaFree(devPtrC);
}
```

```
// Kernel definition, see also section 2.1 of NVIDIA CUDA Programming Guide
__global__ void arrAdd(float *A, float *B, float *C)
{
    // threadIdx.x is a built-in variable provided by CUDA at runtime
    // It represents the thread index inside the block

    int id = threadIdx.x; // id: unique thread identifier

    C[id] = A[id] + B[id];
}
```