

Point-to-point communications

Timing parallel programs

universidade de aveiro



Computação Paralela

Módulo MPI

2021/2022

Rui Costa

Email: americo.costa@ua.pt

Point-to-point communications

Standard send (blocking)

```
MPI_Send(address, count, datatype, destination,  
tag, comm)
```

- `(address, count, datatype)` **describes** count occurrences of items of the form `datatype` starting at `address`,
- `destination` is the *rank* of the destination in the group associated with the *communicator* `comm`,
- `tag` is an integer used for message matching, and
- `comm` is the *communicator*, identifies a group of processes and a communication context.

Point-to-point communications

Standard receive (blocking)



```
MPI_Recv(address, maxcount, datatype, source,  
tag, comm, status)
```

- (address, maxcount, datatype) are the same as in MPI_Send, although it is allowed for less than maxcount occurrences to be received,
- tag and comm are as in MPI_Send, with the addition that a wildcard, matching any tag, is allowed.
- The source is the *rank* of the source of the message in the group associated with the *communicator* comm, or a wildcard matching any source.
- Finally, status holds information about the actual message size, source, and tag, useful when wild cards have been used.

Point-to-point communications

Synchronous send

```
MPI_Ssend(address, count, datatype, destination,  
tag, comm)
```

MPI_Ssend has the same arguments as MPI_Send, but only returns when *receiver* process finishes receiving the message.

The point of the synchronous send operations is avoiding the *sender* process to change the values to be sent before the sending actually occurs.

Point-to-point communications

Buffered send

```
MPI_Bsend(address, count, datatype, destination,  
tag, comm)
```

`MPI_Bsend` has similar arguments as `MPI_Send` and `MPI_Ssend`, but uses a buffer to store the message while the *receiver* is not ready.

This way, the *sender* process can proceed without the risk of overwriting the message to be sent.

In buffered sends, it is necessary need to allocate enough memory for the buffer, and attach/detach it with:

```
MPI_Buffer_attach(buffer, count) ;  
MPI_Buffer_detach(buffer, count) .
```

Timing of parallel programs is especially relevant, since the goal of parallelization is to reduce execution time.

MPI provides a function for timing programs, and sections of programs:

```
MPI_Wtime()
```

Calling `MPI_Wtime()` returns the number of seconds that have passed since some arbitrary point of time in the past, which does not change during the execution of the process.

Elapsed time can be measured with the difference two calls of `MPI_Wtime()`.

The resolution of the output of `MPI_Wtime` is hardware dependent, and can be found by calling

```
MPI_Wtick()
```

Another function that becomes useful for timing programs is

```
MPI_Barrier(comm)
```

This function is a collective operation that does not let the calling process to continue until all processes in the communicator `comm` have called `MPI_Barrier`.

Using MPI: portable parallel programming with the message-passing interface, 3rd edition, William Gropp, Ewing Lusk, and Anthony Skjellum, MIT press (2014).